

Alexa_NLP_analysis

November 18, 2019

1 Exploration of Alexa Reviews using NLP

Pooja Umathe

In this project I am going to predict user given ratings of Alexa using natural language processing techniques. The data consists of the following fields: rating, date, variation, and review text from amazon

```
[1]: # Importing library and dataset
import pandas as pd
df = pd.read_table('amazon_alexa.tsv')
```

C:\Users\pooja\Anaconda3\lib\site-packages\ipykernel_launcher.py:3:

FutureWarning: read_table is deprecated, use read_csv instead, passing sep='\t'.

This is separate from the ipykernel package so we can avoid doing imports until

1.1 Data Preprocessing

```
[2]: # Showing first 5 columns of the data
df.head()
```

```
[2]: rating      date      variation \
0         5  31-Jul-18  Charcoal Fabric
1         5  31-Jul-18  Charcoal Fabric
2         4  31-Jul-18   Walnut Finish
3         5  31-Jul-18  Charcoal Fabric
4         5  31-Jul-18  Charcoal Fabric

                                verified_reviews  feedback
0                                Love my Echo!         1
1                                Loved it!           1
2  Sometimes while playing a game, you can answer...  1
3  I have had a lot of fun with this thing. My 4 ...  1
4                                Music              1
```

```
[3]: # Getting info of the data
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3150 entries, 0 to 3149
Data columns (total 5 columns):
rating                3150 non-null int64
date                  3150 non-null object
variation              3150 non-null object
verified_reviews      3150 non-null object
feedback              3150 non-null int64
dtypes: int64(2), object(3)
memory usage: 123.1+ KB

```

```

[4]: # Describing the data and getting mean, max, min, std, count
df.describe()

```

```

[4]:
count    rating    feedback
count  3150.000000  3150.000000
mean     4.463175    0.918413
std      1.068506    0.273778
min       1.000000    0.000000
25%       4.000000    1.000000
50%       5.000000    1.000000
75%       5.000000    1.000000
max       5.000000    1.000000

```

```

[5]: # Showing columns of the data
df.columns

```

```

[5]: Index(['rating', 'date', 'variation', 'verified_reviews', 'feedback'],
dtype='object')

```

```

[6]: # As we can see that our one column is containing categorical values
# we are using one hot encoder which will encoding for the model to understand
→ this variable

df1 = df[['variation']]
df2 = df.drop(['variation'], axis = 1)
df1 = pd.get_dummies(df1)
df = pd.concat([df1, df2], axis = 1)
df.head()

```

```

[6]:
variation_Black  variation_Black  Dot  variation_Black  Plus  \
0                0                0                0
1                0                0                0
2                0                0                0
3                0                0                0
4                0                0                0

variation_Black  Show  variation_Black  Spot  variation_Charcoal  Fabric  \
0                0                0                0                1
1                0                0                0                1

```

2	0	0	0
3	0	0	1
4	0	0	1

	variation_Configuration: Fire TV Stick	variation_Heather Gray Fabric	\
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

	variation_Oak Finish	variation_Sandstone Fabric	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	variation_Walnut Finish	variation_White	variation_White Dot	\
0	0	0	0	
1	0	0	0	
2	1	0	0	
3	0	0	0	
4	0	0	0	

	variation_White Plus	variation_White Show	variation_White Spot	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	rating	date	verified_reviews	\
0	5	31-Jul-18	Love my Echo!	
1	5	31-Jul-18	Loved it!	
2	4	31-Jul-18	Sometimes while playing a game, you can answer...	
3	5	31-Jul-18	I have had a lot of fun with this thing. My 4 ...	
4	5	31-Jul-18	Music	

	feedback
0	1
1	1
2	1
3	1
4	1

```
[7]: # We have date column in our dataset which needs to be addressed and the model
      →which will accept only numeric inputs,
      # will not be able to interpret this information.
```

```
df['Year'] = pd.DatetimeIndex(df['date']).year
df['Month'] = pd.DatetimeIndex(df['date']).month
df['Day'] = pd.DatetimeIndex(df['date']).day
df = df.drop(['date'], axis = 1)
df.head()
```

```
[7]:
```

	variation_Black	variation_Black	Dot	variation_Black	Plus	\
0	0		0		0	
1	0		0		0	
2	0		0		0	
3	0		0		0	
4	0		0		0	

	variation_Black	Show	variation_Black	Spot	variation_Charcoal	Fabric	\
0		0		0			1
1		0		0			1
2		0		0			0
3		0		0			1
4		0		0			1

	variation_Configuration: Fire TV Stick	variation_Heather	Gray	Fabric	\
0		0			0
1		0			0
2		0			0
3		0			0
4		0			0

	variation_Oak Finish	variation_Sandstone	Fabric	...	\
0	0			0	...
1	0			0	...
2	0			0	...
3	0			0	...
4	0			0	...

	variation_White	Dot	variation_White	Plus	variation_White	Show	\
0		0		0		0	
1		0		0		0	
2		0		0		0	
3		0		0		0	
4		0		0		0	

	variation_White	Spot	rating	\
0		0	5	

```

1          0      5
2          0      4
3          0      5
4          0      5

```

```

                                verified_reviews  feedback  Year  Month  \
0                                Love my Echo!          1  2018    7
1                                Loved it!            1  2018    7
2  Sometimes while playing a game, you can answer...    1  2018    7
3  I have had a lot of fun with this thing. My 4 ...    1  2018    7
4                                Music                1  2018    7

```

```

    Day
0    31
1    31
2    31
3    31
4    31

```

[5 rows x 22 columns]

Now, we are going to use TextBlob. It will analyze text strings and provide numeric outputs relating to sentiment. TextBlob will output both the polarity (range -1 to 1) as well as subjectivity (range 0 to 1) in the form of a tuple. The results of this process can be mapped into columns containing polarity and subjectivity respectively.

```

[8]: from textblob import TextBlob

[9]: # Defining the function which will output sentiment given a string input and
    → that
    # will return a neutral output if the input cannot be handled.

    def sentiment_calc(text):
        try:
            return TextBlob(text).sentiment
        except:
            return TextBlob('hello').sentiment

[10]: # Creating new column sentiment

df['sentiment'] = df['verified_reviews'].apply(lambda text:
    → sentiment_calc(text))
df[['polarity', 'subjectivity']] = df['sentiment'].apply(pd.Series)
df = df.drop(['sentiment'], axis = 1)
df.head()

```

```

[10]: variation_Black  variation_Black  Dot  variation_Black  Plus  \
0                    0                    0                    0
1                    0                    0                    0

```

2	0	0	0
3	0	0	0
4	0	0	0

	variation_Black	Show	variation_Black	Spot	variation_Charcoal	Fabric	\
0		0		0		1	
1		0		0		1	
2		0		0		0	
3		0		0		1	
4		0		0		1	

	variation_Configuration: Fire TV Stick	variation_Heather	Gray	Fabric	\
0		0		0	
1		0		0	
2		0		0	
3		0		0	
4		0		0	

	variation_Oak Finish	variation_Sandstone	Fabric	...	\
0	0		0	...	
1	0		0	...	
2	0		0	...	
3	0		0	...	
4	0		0	...	

	variation_White	Show	variation_White	Spot	rating	\
0		0		0	5	
1		0		0	5	
2		0		0	4	
3		0		0	5	
4		0		0	5	

	verified_reviews	feedback	Year	Month	\
0	Love my Echo!	1	2018	7	
1	Loved it!	1	2018	7	
2	Sometimes while playing a game, you can answer...	1	2018	7	
3	I have had a lot of fun with this thing. My 4 ...	1	2018	7	
4	Music	1	2018	7	

	Day	polarity	subjectivity
0	31	0.625	0.6000
1	31	0.875	0.8000
2	31	-0.100	0.5125
3	31	0.350	0.4500
4	31	0.000	0.0000

[5 rows x 24 columns]

```
[11]: # To return the average word length by defining a function
```

```
def avg_word(sentence):
    try:
        words = sentence.split()
        return (sum(len(word) for word in words)/len(words))
    except:
        return 0
```

```
[12]: # Based on the review text column creating new columns
```

```
df['number_words'] = df['verified_reviews'].str.split().str.len()
df['number_character'] = df['verified_reviews'].str.len()
df['avg_word'] = df['verified_reviews'].apply(lambda x: avg_word(x))
```

```
[13]: # creating a column that contains the number of stopwords in the review text.
```

```
import nltk
nltk.download("stopwords")
from nltk.corpus import stopwords
stop = stopwords.words('english')
df['stopwords'] = df['verified_reviews'].apply(lambda x: len([x for x in x.
    ↳split() if x in stop]))
df.head()
```

[nltk_data] Downloading package stopwords to

[nltk_data] C:\Users\pooja\AppData\Roaming\nltk_data...

[nltk_data] Package stopwords is already up-to-date!

```
[13]:
```

	variation_Black	variation_Black	Dot	variation_Black	Plus	\
0	0		0		0	
1	0		0		0	
2	0		0		0	
3	0		0		0	
4	0		0		0	

	variation_Black	Show	variation_Black	Spot	variation_Charcoal	Fabric	\
0		0		0			1
1		0		0			1
2		0		0			0
3		0		0			1
4		0		0			1

	variation_Configuration: Fire TV Stick	variation_Heather Gray Fabric	\
0	0		0
1	0		0
2	0		0
3	0		0

4			0			0	
	variation_Oak	Finish	variation_Sandstone	Fabric	...	feedback	Year \
0		0		0	...	1	2018
1		0		0	...	1	2018
2		0		0	...	1	2018
3		0		0	...	1	2018
4		0		0	...	1	2018

	Month	Day	polarity	subjectivity	number_words	number_character	\
0	7	31	0.625	0.6000	3	13	
1	7	31	0.875	0.8000	2	9	
2	7	31	-0.100	0.5125	38	195	
3	7	31	0.350	0.4500	34	172	
4	7	31	0.000	0.0000	1	5	

	avg_word	stopwords
0	3.666667	1
1	4.000000	0
2	4.131579	19
3	4.088235	12
4	5.000000	0

[5 rows x 28 columns]

[14]: *# Using Minmaxscaler for better inetrpretation*

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
columns = ['number_character', 'number_words', 'avg_word', 'stopwords']
for col in columns:
    df[[col]] = scaler.fit_transform(df[[col]])
df.head()
```

[14]:	variation_Black	variation_Black	Dot	variation_Black	Plus	\
0		0	0		0	
1		0	0		0	
2		0	0		0	
3		0	0		0	
4		0	0		0	

	variation_Black	Show	variation_Black	Spot	variation_Charcoal	Fabric	\
0		0		0		1	
1		0		0		1	
2		0		0		0	
3		0		0		1	
4		0		0		1	

	variation_Configuration: Fire TV Stick	variation_Heather Gray Fabric	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	variation_Oak Finish	variation_Sandstone Fabric	...	feedback	Year	\
0	0	0	...	1	2018	
1	0	0	...	1	2018	
2	0	0	...	1	2018	
3	0	0	...	1	2018	
4	0	0	...	1	2018	

	Month	Day	polarity	subjectivity	number_words	number_character	\
0	7	31	0.625	0.6000	0.005703	0.004211	
1	7	31	0.875	0.8000	0.003802	0.002807	
2	7	31	-0.100	0.5125	0.072243	0.068070	
3	7	31	0.350	0.4500	0.064639	0.060000	
4	7	31	0.000	0.0000	0.001901	0.001404	

	avg_word	stopwords
0	0.056410	0.004484
1	0.061538	0.000000
2	0.063563	0.085202
3	0.062896	0.053812
4	0.076923	0.000000

[5 rows x 28 columns]

```
[15]: from gensim.summarization import summarize
def sum_text(text):
    try:
        summed_text = summarize(text)
        return summed_text
    except:
        return text
```

```
[16]: df['summed_text'] = df['verified_reviews'].apply(lambda x: sum_text(x))
df['sentiment_sum'] = df['summed_text'].apply(lambda text: sentiment_calc(text))
df[['polarity_sum', 'subjectivity_sum']] = df['sentiment_sum'].apply(pd.Series)
df = df.drop(['sentiment_sum'], axis = 1)
df.head()
```

	variation_Black	variation_Black	Dot	variation_Black	Plus	\
0	0	0	0	0		
1	0	0	0	0		
2	0	0	0	0		

3	0	0	0
4	0	0	0

	variation_Black	Show	variation_Black	Spot	variation_Charcoal	Fabric	\
0		0		0		1	
1		0		0		1	
2		0		0		0	
3		0		0		1	
4		0		0		1	

	variation_Configuration: Fire TV Stick	variation_Heather	Gray	Fabric	\
0		0		0	
1		0		0	
2		0		0	
3		0		0	
4		0		0	

	variation_Oak	Finish	variation_Sandstone	Fabric	...	Day	polarity	\
0		0		0	...	31	0.625	
1		0		0	...	31	0.875	
2		0		0	...	31	-0.100	
3		0		0	...	31	0.350	
4		0		0	...	31	0.000	

	subjectivity	number_words	number_character	avg_word	stopwords	\
0	0.6000	0.005703	0.004211	0.056410	0.004484	
1	0.8000	0.003802	0.002807	0.061538	0.000000	
2	0.5125	0.072243	0.068070	0.063563	0.085202	
3	0.4500	0.064639	0.060000	0.062896	0.053812	
4	0.0000	0.001901	0.001404	0.076923	0.000000	

	summed_text	polarity_sum	subjectivity_sum
0	Love my Echo!	0.625	0.6
1	Loved it!	0.875	0.8
2		0.000	0.0
3		0.000	0.0
4	Music	0.000	0.0

[5 rows x 31 columns]

```
[17]: # Creating a new dataframe which can be fed into the model
```

```
df2 = df.drop(['verified_reviews', 'summed_text'], axis = 1)
df2.head()
```

```
[17]:
```

	variation_Black	variation_Black	Dot	variation_Black	Plus	\
0	0		0		0	
1	0		0		0	

2	0	0	0
3	0	0	0
4	0	0	0

	variation_Black	Show	variation_Black	Spot	variation_Charcoal	Fabric	\
0		0		0		1	
1		0		0		1	
2		0		0		0	
3		0		0		1	
4		0		0		1	

	variation_Configuration: Fire TV Stick	variation_Heather	Gray	Fabric	\
0		0		0	
1		0		0	
2		0		0	
3		0		0	
4		0		0	

	variation_Oak	Finish	variation_Sandstone	Fabric	...	Month	Day	\
0		0		0	...	7	31	
1		0		0	...	7	31	
2		0		0	...	7	31	
3		0		0	...	7	31	
4		0		0	...	7	31	

	polarity	subjectivity	number_words	number_character	avg_word	\
0	0.625	0.6000	0.005703	0.004211	0.056410	
1	0.875	0.8000	0.003802	0.002807	0.061538	
2	-0.100	0.5125	0.072243	0.068070	0.063563	
3	0.350	0.4500	0.064639	0.060000	0.062896	
4	0.000	0.0000	0.001901	0.001404	0.076923	

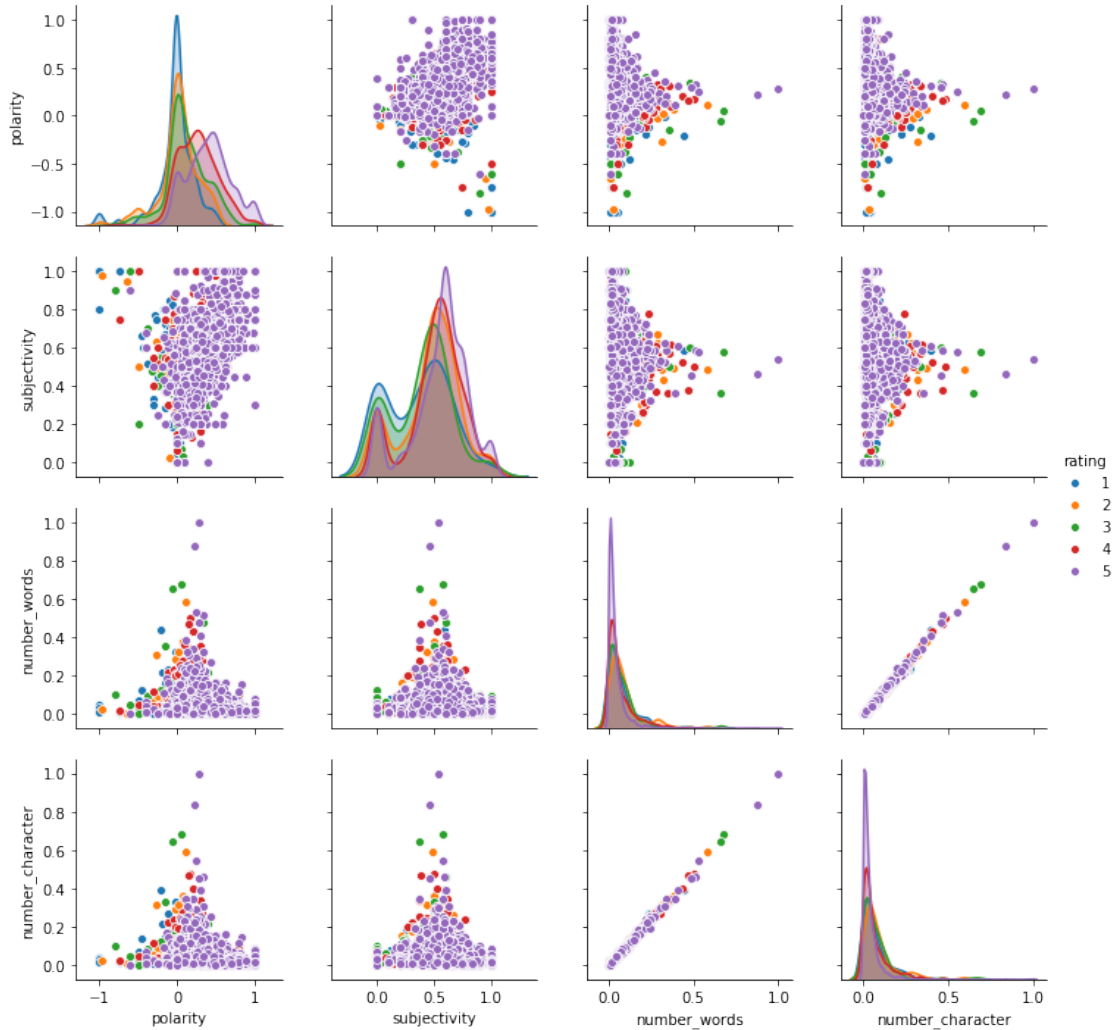
	stopwords	polarity_sum	subjectivity_sum
0	0.004484	0.625	0.6
1	0.000000	0.875	0.8
2	0.085202	0.000	0.0
3	0.053812	0.000	0.0
4	0.000000	0.000	0.0

[5 rows x 29 columns]

1.2 Data Visualization

```
[18]: import seaborn as sns
      from matplotlib import pyplot as plt
```

```
[19]: sns_plot = sns.pairplot(df, hue = 'rating', vars=['polarity', 'subjectivity', 'number_words', 'number_character'])
#sns_plot.savefig('name_of_file.png')
#files.download('name_of_file.png')
```



1.3 Building and Executing the Model

```
[20]: from sklearn.model_selection import train_test_split
import numpy as np
df1=np.matrix(df2.drop(['rating'], axis = 1))
y=np.array(df2['rating'])
X_train, X_test, y_train, y_test = train_test_split(df1, y, test_size=0.2)
```

```
[21]: from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
```

```

rfc = RandomForestClassifier(n_estimators=500, criterion='gini', max_depth=None,
                             min_samples_split=2, min_samples_leaf=1,
                             →min_weight_fraction_leaf=0.0,
                             max_features='auto', max_leaf_nodes=None,
                             →min_impurity_decrease=0.0,
                             min_impurity_split=None, bootstrap=True,
                             →oob_score=False, n_jobs=None,
                             random_state=None, verbose=0, warm_start=False,
                             →class_weight=None)

rfc.fit(X_train,y_train)

score = rfc.score(X_test,y_test)

print(score)

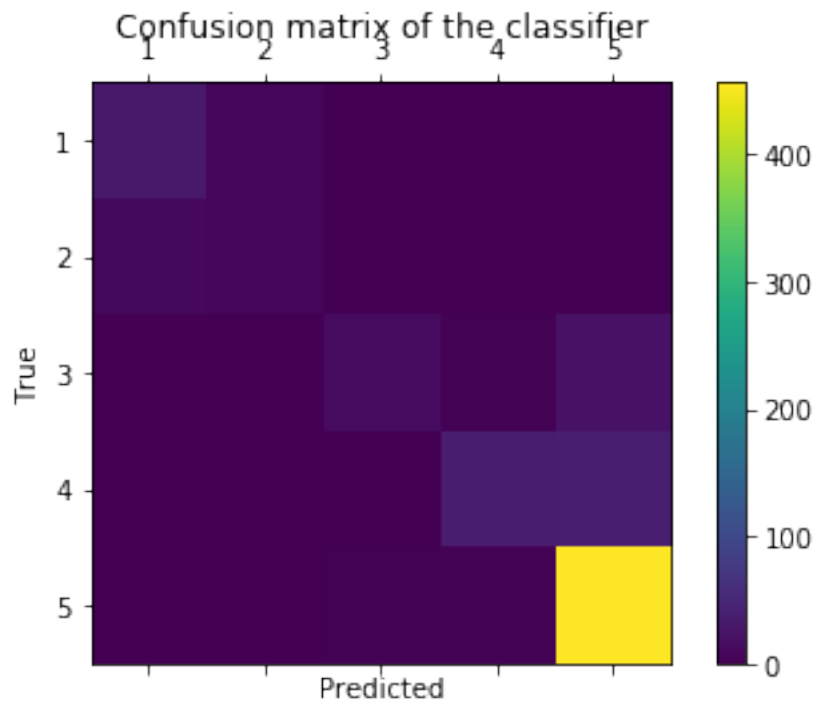
```

0.8634920634920635

```

[22]: from sklearn.metrics import confusion_matrix
predictions = rfc.predict(X_test)
cm = confusion_matrix(y_test, predictions)
labels = ['1','2','3','4','5']
fig = plt.figure()
ax = fig.add_subplot(111)
cax = ax.matshow(cm)
plt.title('Confusion matrix of the classifier')
fig.colorbar(cax)
ax.set_xticklabels([''] + labels)
ax.set_yticklabels([''] + labels)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()

```



[23]: `print(cm)`

```
[[ 31   6   0   0   0]
 [ 10   8   0   0   0]
 [  0   0  13   3  22]
 [  0   0   0  35  39]
 [  0   0   3   3 457]]
```

The results are pretty good. We achieved an accuracy of approximately 84%, with most of the incorrect predictions being in an adjacent square.